

Modeling_And_Evaluation

October 11, 2024

0.0.1 Data Modeling

```
[3]: #importing the necessary libraries
import pandas as pd
import numpy as np
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix, roc_auc_score, log_loss
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import LabelBinarizer
from xgboost import XGBClassifier
```

```
[4]: #loading the preprocessed data
X_train = pd.read_csv('./data/X_train.csv')
X_test = pd.read_csv('./data/X_test.csv')
y_train = pd.read_csv('./data/y_train.csv')
y_test = pd.read_csv('./data/y_test.csv')

y_train = y_train.values.ravel()
y_test = y_test.values.ravel()
```

```
[5]: #print the data shapes of the training and testing data
print("Data loaded successfully.")
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")
```

```
Data loaded successfully.
X_train shape: (299, 22)
X_test shape: (75, 22)
y_train shape: (299,)
y_test shape: (75,)
```

```
[6]: #define model evaluation function
def evaluate_model(model, X, y, model_name):
    cv_accuracy = cross_val_score(model, X, y, cv=5, scoring='accuracy')
    cv_log_loss = cross_val_score(model, X, y, cv=5, scoring='neg_log_loss')

    print(f"\n{model_name} Cross-Validation Results:")
    print(f"Mean Accuracy: {cv_accuracy.mean():.4f} (+/- {cv_accuracy.std() * 2:
↪.4f})")
    print(f"Mean Log Loss: {-cv_log_loss.mean():.4f} (+/- {cv_log_loss.std() * 2:
↪.4f})")
```

```
[7]: #define model evaluation function
def train_predict_evaluate(model, X_train, X_test, y_train, y_test, model_name):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    y_pred_proba = model.predict_proba(X_test)

    accuracy = accuracy_score(y_test, y_pred)

    lb = LabelBinarizer()
    y_test_bin = lb.fit_transform(y_test)
    roc_auc = roc_auc_score(y_test_bin, y_pred_proba, multi_class='ovr',
↪average='macro')

    logloss = log_loss(y_test, y_pred_proba)

    print(f"\n{model_name} Test Results:")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"ROC AUC: {roc_auc:.4f}")
    print(f"Log Loss: {logloss:.4f}")
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred))
    print("\nConfusion Matrix:")
    print(confusion_matrix(y_test, y_pred))

    return {
        'model_name': model_name,
        'accuracy': accuracy,
        'roc_auc': roc_auc,
        'log_loss': logloss
    }
```

```
[22]: #list of models to evaluate
models = [
    (LogisticRegression(multi_class='ovr', random_state=42), "Multinomial
↪Logistic Regression"),
    (DecisionTreeClassifier(random_state=42), "Decision Tree"),
```

```

(RandomForestClassifier(n_estimators=100, random_state=42), "Random_
↪Forest"),
(XGBClassifier(random_state=42), "XGBoost"),
(SVC(kernel='rbf', random_state=42, probability=True), "Support Vector_
↪Machine"),
(KNeighborsClassifier(n_neighbors=5), "K-Nearest Neighbors"),
]

results = []

for model, name in models:
    print(f"\n{'='*50}\nEvaluating {name}\n{'='*50}")
    evaluate_model(model, X_train, y_train, name)
    result = train_predict_evaluate(model, X_train, X_test, y_train, y_test,
↪name)
    results.append(result)

```

```

=====
Evaluating Multinomial Logistic Regression
=====

```

Multinomial Logistic Regression Cross-Validation Results:

Mean Accuracy: 0.9131 (+/- 0.0827)

Mean Log Loss: 0.4210 (+/- 0.3257)

Multinomial Logistic Regression Test Results:

Accuracy: 0.9067

ROC AUC: 0.9221

Log Loss: 0.3854

Classification Report:

	precision	recall	f1-score	support
0	0.78	0.88	0.82	16
1	0.95	0.95	0.95	43
2	0.93	0.81	0.87	16
accuracy			0.91	75
macro avg	0.89	0.88	0.88	75
weighted avg	0.91	0.91	0.91	75

Confusion Matrix:

```

[[14  1  1]
 [ 2 41  0]

```

[2 1 13]]

=====

Evaluating Decision Tree

=====

Decision Tree Cross-Validation Results:

Mean Accuracy: 0.8763 (+/- 0.1255)

Mean Log Loss: 3.2240 (+/- 2.6318)

Decision Tree Test Results:

Accuracy: 0.8933

ROC AUC: 0.8873

Log Loss: 3.9004

Classification Report:

	precision	recall	f1-score	support
0	0.76	0.81	0.79	16
1	0.95	0.98	0.97	43
2	0.86	0.75	0.80	16
accuracy			0.89	75
macro avg	0.86	0.85	0.85	75
weighted avg	0.89	0.89	0.89	75

Confusion Matrix:

[[13 1 2]
[1 42 0]
[3 1 12]]

=====

Evaluating Random Forest

=====

Random Forest Cross-Validation Results:

Mean Accuracy: 0.9064 (+/- 0.0831)

Mean Log Loss: 2.0999 (+/- 3.3583)

Random Forest Test Results:

Accuracy: 0.8800

ROC AUC: 0.9208

Log Loss: 2.1108

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.72	0.81	0.76	16
1	0.95	0.98	0.97	43
2	0.85	0.69	0.76	16
accuracy			0.88	75
macro avg	0.84	0.83	0.83	75
weighted avg	0.88	0.88	0.88	75

Confusion Matrix:

```
[[13  1  2]
 [ 1 42  0]
 [ 4  1 11]]
```

=====

Evaluating XGBoost

=====

XGBoost Cross-Validation Results:

Mean Accuracy: 0.9031 (+/- 0.0826)

Mean Log Loss: 0.6215 (+/- 0.5153)

XGBoost Test Results:

Accuracy: 0.9067

ROC AUC: 0.9011

Log Loss: 0.6244

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.81	0.81	16
1	0.95	0.98	0.97	43
2	0.87	0.81	0.84	16
accuracy			0.91	75
macro avg	0.88	0.87	0.87	75
weighted avg	0.91	0.91	0.91	75

Confusion Matrix:

```
[[13  1  2]
 [ 1 42  0]
 [ 2  1 13]]
```

=====

Evaluating Support Vector Machine

=====

Support Vector Machine Cross-Validation Results:

Mean Accuracy: 0.9030 (+/- 0.0929)

Mean Log Loss: 0.3505 (+/- 0.2622)

Support Vector Machine Test Results:

Accuracy: 0.8800

ROC AUC: 0.9108

Log Loss: 0.3997

Classification Report:

	precision	recall	f1-score	support
0	0.72	0.81	0.76	16
1	0.95	0.98	0.97	43
2	0.85	0.69	0.76	16
accuracy			0.88	75
macro avg	0.84	0.83	0.83	75
weighted avg	0.88	0.88	0.88	75

Confusion Matrix:

```
[[13  1  2]
 [ 1 42  0]
 [ 4  1 11]]
```

=====

Evaluating K-Nearest Neighbors

=====

K-Nearest Neighbors Cross-Validation Results:

Mean Accuracy: 0.8897 (+/- 0.0883)

Mean Log Loss: 2.2124 (+/- 3.2205)

K-Nearest Neighbors Test Results:

Accuracy: 0.8800

ROC AUC: 0.9048

Log Loss: 3.0221

Classification Report:

	precision	recall	f1-score	support
0	0.72	0.81	0.76	16
1	0.95	0.98	0.97	43
2	0.85	0.69	0.76	16
accuracy			0.88	75
macro avg	0.84	0.83	0.83	75

weighted avg	0.88	0.88	0.88	75
--------------	------	------	------	----

Confusion Matrix:

```
[[13  1  2]
 [ 1 42  0]
 [ 4  1 11]]
```

```
[23]: #Compare the results of the models
comparison_df = pd.DataFrame(results)
comparison_df = comparison_df.set_index('model_name')
print("\nModel Comparison:")
display(comparison_df)

best_accuracy = comparison_df['accuracy'].idxmax()
best_roc_auc = comparison_df['roc_auc'].idxmax()
best_log_loss = comparison_df['log_loss'].idxmin()

print(f"\nBest model by accuracy: {best_accuracy}")
print(f"Best model by ROC AUC: {best_roc_auc}")
print(f"Best model by log loss: {best_log_loss}")
```

Model Comparison:

	accuracy	roc_auc	log_loss
model_name			
Multinomial Logistic Regression	0.906667	0.922144	0.385404
Decision Tree	0.893333	0.887287	3.900374
Random Forest	0.880000	0.920789	2.110767
XGBoost	0.906667	0.901072	0.624433
Support Vector Machine	0.880000	0.910795	0.399662
K-Nearest Neighbors	0.880000	0.904770	3.022085

Best model by accuracy: Multinomial Logistic Regression

Best model by ROC AUC: Multinomial Logistic Regression

Best model by log loss: Multinomial Logistic Regression